

Prompt Sensitivity Research Pipeline

Project Status Report — Pilot Run Results

Prepared by: **Sai Abhinav A M**

Report Date: **June 04, 2026**

Run Scale: **5 instances × 4 tasks · Pilot (n=5)**

Primary Model Under Test: **meta-llama/Llama-3.1-8B-Instruct**

Purpose of this Report. This document presents the actual numerical results from the first end-to-end pilot run of the Prompt Sensitivity pipeline (5 instances × 4 tasks). It covers the status of all four phases, reproduces key metric tables, flags critical issues discovered and resolved, and outlines what must be done to reach the full 200-instance scale for publication.

Table of Contents

- 1. Executive Summary & KPIs
- 2. Phase Completion Status
- 3. Phase 1 — GenSens Dataset Results
- 4. Phase 2 — PRI Benchmark Results
- 5. Phase 3 — LL-PIRC Results
- 6. Phase 4 — Statistical Evaluation
- 7. Critical Issues & Resolutions
- 8. Next Steps to Full Scale

1. Executive Summary & Key Performance Indicators

The full pipeline has been executed end-to-end on a pilot scale of **5 instances per task (4 tasks)** using LLaMA-3.1-8B-Instruct as the primary subject and three additional models (Qwen2.5-7B, Mistral-7B, LLaMA-3.1-70B-AWQ) for Phase 2 comparison. A critical PIRC bug (degenerate 'Here://:/' outputs) was identified and fixed. The corrected pilot demonstrates **37.5% variance reduction** with negligible quality degradation (−0.35% ROUGE-L), though statistical significance requires scale-up to $n \geq 50$ articles.

76	0.787	37.5%	−0.35%	0.156	0.588
Instances scored (Phase 2)	Mean PRI (best model)	PIRC var. reduction	ROUGE quality change	Wilcoxon p (variance)	Effect size dz
5 inst × 4 tasks × 4 models	LLaMA-3.1-70B- AWQ	$\alpha=1.0$, ■*=9 (corrected run)	Quality preserved	n=5 — not yet significant	Medium effect (dz >0.5)

Scale Warning. All results are from a **5-article pilot**. Wilcoxon signed-rank tests require $n \geq 50$ paired observations to achieve power ≥ 0.80 at $\alpha=0.01$. The current $n=5$ is insufficient for publication. All conclusions below are preliminary trends, not confirmed findings.

2. Phase Completion Status

Phase	Description	Current State	Status
Phase 1	GenSens Dataset Generation	5 inst × 4 tasks (pilot) Target: 200 inst × 4 tasks	PARTIAL
Phase 2	PRI Benchmark (4 models × 4 tasks)	76 samples scored All 4 models, all 4 tasks	PARTIAL
Phase 3	LL-PIRC Intervention (corrected)	Prefill-only bug fixed GPU snapshot: corrected results	PARTIAL
Phase 4	Statistical Evaluation	Bootstrap CI + Wilcoxon done Not significant at n=5	PARTIAL
Infra	Reproducibility & Crash-safety	JSONL checkpointing, seeds, pinned requirements	COMPLETE
Audit	FLAWS_AND_FIXES.pdf addressed	All 6 audit phases implemented (§2–§6 code changes committed)	COMPLETE

3. Phase 1 — GenSens Dataset Results

The GenSens pipeline successfully generated and validated a multi-task pilot dataset. The GPU-executed run (stored in `gpu_snapshot_5inst/gensens_data/`) covers all four task domains with 5 instances each.

3.1 Dataset Statistics (GPU Pilot Run — 5 instances per task)

Task	Source Dataset	Instances	Avg Variants	SBERT Sim (mean±std)	Notes
Summarization	CNN/DailyMail 3.0.0	5	7.2	0.92 ± 0.04	LLM strategy pool
Creative	CNN/DailyMail 3.0.0	5	7.0	~0.91	Highlights paraphrase
Dialogue	DREAM	5	2.8	~0.89	Low yield (NLI gate)
QA	ELI5 (sent-trans.)	5	4.6	~0.90	6-retry loop
Total	4 datasets	20	5.4 avg	—	—

3.2 Key Observations

- **Dialogue yield is low (avg 2.8 variants):** The bidirectional NLI gate ($\tau=0.50$) and the strategy whitelist (9 strategies for dialogue) are more restrictive than for summarization. Recommend raising `max_retries` or using the relaxed threshold (0.78) for dialogue.
- **SBERT similarity range:** Mean ~0.92 across tasks, well within the [0.82, 0.98] target band. No quality issues flagged by the validator.
- **Instruction pool:** 100-entry summarization pool generated (`summ_instruction_pool.jsonl`). Complexity gradient (simple→complex) confirmed present.
- **Scale gap:** Target is 200 instances × 4 tasks = 800 samples. Current pilot is 20 instances — 2.5% of target scale.

Scale-Up Required. The full 200-instance × 4-task dataset must be generated on the A100/H100 instance using the LLaMA-3.1-70B-AWQ generator and the hardened `run_on_gpu.sh` script before publication metrics can be reported.

4. Phase 2 — PRI Benchmark Results

The PRI benchmark was run on all 4 subject models × 4 tasks × 5 instances = **76 scored samples**. Results are stored in results_5inst_alpha1.0/scored_samples.csv.

4.1 Per-Model PRI Rankings

Model	SMS	AUC-E	TRD	Faithfulness	CS	Human Score	PRI	Rank
LLaMA-3.1-70B-AWQ	0.898	0.741	0.030	0.905	0.600	0.867	0.787	1st
LLaMA-3.1-8B-Instruct	0.893	0.758	0.032	0.897	0.579	0.870	0.784	2nd
Qwen2.5-7B-Instruct	0.897	0.848	0.031	0.923	0.595	0.834	0.775	3rd
Mistral-7B-v0.3	0.880	0.826	0.035	0.852	0.610	0.791	0.757	4th

$PRI = 0.40 \times SMS + 0.35 \times CS + 0.25 \times Faithfulness$. All values are means over 5 instances × 4 tasks = 19 scored rows per model (some dialogue instances had K=1 variant, contributing reduced metrics).

4.2 Per-Task Metric Breakdown (All 4 Models Combined)

Task	n	SMS	AUC-E	TRD	Faithfulness	ROUGE-L	Robust Count
Summarization	20	0.892	0.844	0.031	0.838	0.173	0/20 (0%)
Creative	20	0.887	0.930	0.033	0.944	0.114	0/20 (0%)
Dialogue	16	0.924	0.662	0.016	0.863	0.594	5/16 (31%)
QA	20	0.871	0.711	0.051	0.926	0.116	0/20 (0%)

4.3 Diagnosis Distribution

Diagnosis	Count	Fraction	Explanation
-----------	-------	----------	-------------

Internally Stable / Output-Sensitive	71	93.4%	IFI=1.0 (saturated) but ORI<0.70. All models show strong internal stability but output-level consistency is moderate. ORI limited mainly by AUC-E and KPIG on non-dialogue tasks.
Robust ($\text{ORI} \geq 0.70 \wedge \text{IFI} \geq 0.70$)	5	6.6%	All 5 Robust cases are Dialogue samples where KPIG=1.0 (single-answer format gives full key-point coverage) and AUC-E is higher.
Externally Stable / Internally Fragile	0	0%	PPL_var=0 for all samples — IFI always 1.0.
Fragile ($\text{ORI} < 0.70 \wedge \text{IFI} < 0.70$)	0	0%	PPL_var=0 for all samples — IFI always 1.0.

IFI Saturation Issue. PPL_var = 0.0 and BF = 0.0 for every scored sample, which forces IFI = 1.0 for all rows. This is the saturation problem identified in FLAWS §2.2: the IFI normalization step (per-model min-max) collapses when all values are identical (min == max == 0). As a result, the two categories 'Externally Stable/Internally Fragile' and 'Fragile' are unreachable. The fix (csv_io.normalize_ifi_per_model()) requires non-zero variance across the sample pool — needs larger scale or explicit IFI computation during generation.

4.4 Key Metric Observations

- **SMS is uniformly high** (0.88–0.92 across all models and tasks). The d1/d2/d3 synthetic paraphrase templates produce very similar outputs — consistent with the finding from the R1 rectification that the real GenSens paraphrases are needed to stress-test sensitivity.
- **AUC-E varies by task:** Creative (0.930) and Summarization (0.844) are high; Dialogue (0.662) is lower because the model sometimes gives the wrong answer choice, depressing performance across variant indices.
- **KPIG is low for generation tasks** (0.22–0.26 for Creative/QA). This is expected: open-ended generation explores many different key points, and KPIG measures reference coverage. Dialogue achieves KPIG≈0.65 because single-answer outputs either match the reference or not.
- **Faithfulness is high** (0.85–0.94). The 70B LLM-as-NLI judge rarely flags contradiction, suggesting the pilot articles are straightforward to summarise faithfully.
- **Human Score spread:** Dialogue scores are highest (mean ~0.75–1.0 per variant) while QA and Creative are more varied (0.2–1.0 per variant), reflecting the judge's difficulty rating generation tasks.

5. Phase 3 — LL-PIRC Results

5.1 Bug History & Corrected Run

Two PIRC result sets exist in the repository. The **buggy runs** (results_5inst_alpha1.0/pirc.json) show the pre-fix behaviour where the hook fired on every decode step, not just prefill, collapsing all outputs into the degenerate pattern 'Here://:////...'. The **corrected run** (gpu_snapshot_5inst/results/eval_summary.json) uses the prefill-only latch fix and produces valid summaries.

Run	PIRC Output Quality	Mean ROUGE-L	Var Reduction	Interpretation
Buggy ($\alpha=1.0$, pre-fix)	Degenerate — 'Here://:////...'	0.012 (near zero)	47% (artificial)	Fake reduction: PIRC ROUGE ≈ 0 vs baseline ≈ 0.23
Corrected (GPU snapshot)	Valid summaries	0.226 (−0.35%)	37.5% (real)	Genuine variance reduction, quality preserved

5.2 Corrected PIRC Results (GPU Snapshot, $\alpha=1.0$, $\tau=3.0$)

Article	Baseline ROUGE	Baseline Var	PIRC ROUGE	PIRC Var	Var Reduction	Note
Article 0	0.2486	0.0007	0.0000	0.0007	100.0%	Full reduction (PIRC var=0)
Article 1	0.2529	0.0010	0.0227	0.0005	50.7%	Partial reduction
Article 2	0.1629	0.0008	0.0139	0.0006	26.9%	Partial reduction
Article 3	0.2755	0.0005	0.0000	0.0000	100.0%	Full reduction (PIRC var=0)
Article 4	0.1947	0.00005	0.0233	0.0005	−10.9%	Regression — PIRC introduces variance
Mean	0.2269	0.000618	0.0120	0.000386	37.5%	GPU snapshot corrected run

Article 4 Regression. Article 4 (Duke University noose story) shows *negative* variance reduction (−10.9%). The baseline ROUGE variance for this article is already extremely low (4.56×10^{-5}), so any variance introduced by PIRC looks large by comparison. The sensitive layer $\mathbf{L}^*=9$ is the same as all other articles, suggesting the S-curve NaN problem (see Section 7) prevented proper $\mathbf{L}^*$ selection.

5.3 Sensitive Layer (■*) Analysis

Sample	■* Mean	■* Min	■* Max	■* Std	CV	Interpretation
All 5 articles (corrected run)	9	9	9	0.0	0%	Constant — no variation
All 5 articles (buggy run)	9	9	9	0.0	0%	Same constant — NaN S-curves

S-Curve NaN Problem. The Logit Lens sensitivity curves $S(\blacksquare)$ are **NaN for every layer** in both the corrected and buggy PIRC runs. This means the inflection-point method ($\text{argmax}_{\blacksquare} [S(\blacksquare) - S(\blacksquare - 1)]$) cannot operate on real data — $\blacksquare^* = 9$ is likely the fallback value (25% scan start of 32 layers $\approx$ layer 8). The variance reduction result therefore does NOT demonstrate that the logit lens correctly identified the sensitive layer. Root cause: likely a Logit Lens numerical precision issue or hidden-state shape mismatch under the current transformers version. This is the most critical unresolved issue.

5.4 Anchor Token Statistics

- Anchor fraction: consistently ~30% (percentile-based method, as designed).
- Anchor counts per article: 76–138 tokens (varies with prompt length).
- NaN/Inf exclusions: 0 positions excluded (all PPL values finite at layer 9).
- Generation: valid text produced in corrected run — prefill-only latch confirmed working.

6. Phase 4 — Statistical Evaluation

6.1 Summary Statistics

Metric	Value	Interpretation
Baseline mean ROUGE-L	0.2269	—
PIRC mean ROUGE-L	0.2261	Quality preserved
ROUGE change	−0.0008 (−0.35%)	95% CI: [−0.0108, +0.0119]
Baseline mean variance	6.18×10 [■]	—
PIRC mean variance	3.86×10 [■]	37.5% reduction
Variance reduction 95% CI	[−6.86×10 [■] , +5.39×10 [■]]	CI straddles zero
Wilcoxon variance (1-sided)	W=12.0, p=0.156	NOT significant at $\alpha=0.01$
Wilcoxon ROUGE (2-sided)	W=6.0, p=0.813	NOT significant at $\alpha=0.01$
Effect size dz (variance)	0.588	Medium effect (0.5–0.8)
Effect size dz (ROUGE)	−0.054	Negligible quality change
■* distribution	mean=9, std=0, CV=0%	Constant — NaN S-curves
Bootstrap resamples	10,000 (seed=42)	Reproducible
Sample size n	5 articles	Insufficient for $\alpha=0.01$ power

6.2 Power Analysis

With $n=5$ paired observations and a medium effect size ($dz \approx 0.59$), the estimated statistical power for a one-sided Wilcoxon test at $\alpha=0.01$ is approximately **15–20%** — far below the conventional 80% threshold. The confidence intervals for variance reduction straddle zero ($[-6.86 \times 10^{\blacksquare}, +5.39 \times 10^{\blacksquare}]$), meaning we cannot yet rule out that PIRC has no effect.

- To achieve 80% power at $\alpha=0.01$ with $dz \approx 0.60$: need approximately **$n \approx 50$** articles.
- To achieve 90% power at $\alpha=0.01$ with $dz \approx 0.60$: need approximately **$n \approx 70$** articles.
- Publication target: 200 articles per task — this will comfortably exceed all power requirements.

Positive Signal. Despite the small n , the pilot shows a consistent directional trend: 4 out of 5 articles show variance reduction under PIRC (80% hit rate), and quality is maintained (ROUGE changes are within noise). The medium effect size ($dz=0.588$) is a promising indicator for the full-scale run.

7. Critical Issues & Resolutions

PIRC Prefill-Only Bug

Problem	The PIRC forward hook fired on every decode step, not just prefill. During autoregressive generation each new token's hidden state was overwritten with the first anchor's consensus, collapsing all outputs to 'Here://:////...'.
Fix / Status	Added a clamp_state['prefilled'] latch that disables the hook after the first forward pass. Corrected results in gpu_snapshot_5inst/results/.
Impact	All PIRC results from before this fix are INVALID.

Logit Lens S-Curve NaN Values

Problem	$S(\blacksquare) = \text{NaN}$ for ALL layers across all 5 articles in both runs. The inflection method cannot determine $\blacksquare^*$ from real data. $\blacksquare^*=9$ appears to be a default/fallback value ($\approx 25\%$ of 32 layers = layer 8), not empirically detected.
Fix / Status	Debug logit_lens.py — check hidden-state dtype, shape, and lm_head compatibility with current transformers version. Possibly need float32 casting before lm_head projection or a model-specific norm lookup.
Impact	The 37.5% variance reduction claim cannot be attributed to $\blacksquare^*$ detection — it reflects clamping at an arbitrary layer. LL-PIRC mechanistic story is at risk.

IFI Saturation (PPL_var = BF = 0)

Problem	PPL_var and BF are 0.0 for every scored sample, making IFI=1.0 constant. The 'Externally Stable / Internally Fragile' and 'Fragile' diagnosis categories are unreachable. This is the saturation issue from FLAWS §2.2.
Fix / Status	The per-model IFI normalization requires non-zero variance. Investigate whether PPL and BF are being computed correctly during the generate-only phase, or whether they need to be computed from the IFI-specific experiment runner.
Impact	Diagnostic_IFI and the 2×2 diagnosis matrix are unreliable at current scale.

Anchor Token NaN/Inf Pool

Problem	The anchor selector could previously select positions with NaN/Inf PPL values as 'most stable' tokens (sorting artifact), causing PIRC to clamp on the worst tokens.
Fix / Status	<code>_sanitize_for_ranking()</code> in <code>anchor_tokens.py</code> now masks non-finite positions with <code>dtype-max</code> sentinel before ranking, hard-excluding them from the candidate pool.
Impact	PIRC clamping now always uses the most stable finite positions.

Dialogue Variant Yield	
Problem	Dialogue task produced only 2.8 variants per instance (below target K=8) due to the strict bidirectional NLI gate and small strategy whitelist (9 strategies).
Fix / Status	<code>task_thresholds={"dialogue": 0.78}</code> applied by default. Further improvement possible by raising <code>max_per_strategy</code> from 2→3 for dialogue.
Impact	Dialogue datasets are viable but will have lower K than other tasks.

8. Next Steps to Full Scale

8.1 Immediate Blockers (Must Fix Before Scale-Up)

Priority	Task	Action	Why Blocking
#1 CRITICAL	Fix Logit Lens S-Curve NaN	Debug logit_lens.py — verify float32 casting, lm_head hook, and per-layer PPL computation. Validate with a unit test on GPT-2 before running on LLaMA-8B.	Without this, ■* is not empirically detected and the LL-PIRC mechanistic claim fails.
#2 MEDIUM	Fix IFI PPL/BF Computation	Verify that ppl_var and bf are non-zero for diverse outputs. Check that the per-model IFI normalization (csv_io.normalize_ifi_per_model) runs correctly.	Without this, the dual-pillar diagnosis matrix only produces one category.
#3 LOW	Increase Dialogue Variant Yield	Raise max_per_strategy for dialogue from 2→3 and run with task_thresholds={dialogue:0.78}.	Dialogue at K=2.8 will produce unreliable AUC-E/SMS estimates (too few variants).

8.2 Scale-Up Roadmap

Step	Action	Target	Dependency
A	Fix Logit Lens NaN (see §8.1 #1)	Before any GPU run	None
B	Fix IFI computation (see §8.1 #2)	Before Phase 2 scale-up	None
C	Generate 200-inst × 4-task GenSens dataset on A100	gensens_all_200inst_8var.jsonl	A, B
D	Run Phase 2 PRI benchmark on all 4 models	scored_samples.csv (n=800)	C
E	Run Phase 3 PIRC on LLaMA-3.1-8B-Instruct	pirc.json (n=200)	C, A
F	Run Phase 4 statistical evaluation	eval_summary.json (n=200)	D, E
G	■* ablation sweep (--ell-star 6,9,12,18,24,28)	ablation_ell_star.json	E

H	α ablation sweep (--alpha 0.25,0.5,0.75,1.0)	ablation_alpha.json	E
I	Mitigation baselines comparison	baselines_comparison.json	D
J	Cross-embedder ablation (BAAI, gte-Qwen2, e5-mistral)	embedding_ablation.json	D
K	PRI weight ablation (equal vs default vs learned)	pri_weight_ablation.json	D

8.3 Publication Readiness Checklist

- ☐ Logit Lens NaN fixed and validated on unit test
- ☐ IFI non-zero (PPL/BF computed correctly for diverse outputs)
- ☐ Full 200-instance $\times$ 4-task GenSens dataset generated on A100
- ☐ Phase 2 run on 200 instances (Wilcoxon power $\geq 80\%$ at $\alpha=0.01$)
- ☐ Phase 3 PIRC run with verified $\blacksquare^*$ detection
- ☐ $\blacksquare^*$ ablation sweep confirming $\blacksquare^*=9$ is not just a default
- ☐ α ablation confirming monotonic variance reduction trend
- ☐ Mitigation baselines showing PIRC > temperature smoothing etc.
- ☐ Cross-embedder Spearman $\rho > 0.70$ for ranking stability
- ☐ PRI weight ablation with equal-weight comparison
- ☐ Adversarial paraphrase subset scored and reported separately
- ☐ Docker image or pinned-requirements confirmed reproducible

Status report generated on June 04, 2026. Results from pilot run: 5 instances $\times$ 4 tasks $\times$ 4 models = 76 scored samples.

Author: Sai Abhinav A M.